

# Virtio网络发包过程分析

📅 2018-08-24 | 📁 2.虚拟化 , virtio | 👁 1596 | 💬 1

📄 19k | ⌚ 18 分钟

## 前言

➕ 本文将分析Virtio网络发包过程，包括Frontend如何发送数据给Backend，以及Backend收到数据后的处理。

## 正文

### Frontend：Transmit packet

查看driver注册的ops函数，其中指定的发送函数为**start\_xmit**，我们从这里开始分析发包过程。

virtio\_net.c : virtnet\_netdev

```
1 static const struct net_device_ops virtnet_netdev = {
2     .ndo_open      = virtnet_open,
3     .ndo_stop      = virtnet_close,
4     .ndo_start_xmit = start_xmit,
5     .ndo_validate_addr = eth_validate_addr,
6     .ndo_set_mac_address = virtnet_set_mac_address,
7     .ndo_set_rx_mode  = virtnet_set_rx_mode,
8     .ndo_change_mtu   = virtnet_change_mtu,
9     .ndo_get_stats64  = virtnet_stats,
10    .ndo_vlan_rx_add_vid = virtnet_vlan_rx_add_vid,
11    .ndo_vlan_rx_kill_vid = virtnet_vlan_rx_kill_vid,
12    .ndo_select_queue   = virtnet_select_queue,
13    #ifdef CONFIG_NET_POLL_CONTROLLER
14    .ndo_poll_controller = virtnet_netpoll,
15    #endif
16    .ndo_features_check  = passthru_features_check,
17 };
```

↑ 90%

**start\_xmit**函数首先获取sk\_buff对应的send\_queue，**free\_old\_xmit\_skbs**释放backend处理过的desc，然后调用**xmit\_skb**发包，之后调用**skb\_orphan**孤立skb,因为对于guest来说**xmit\_skb**之后skb的内容已经是发送的了，最后调用**virtqueue\_kick**通知Host。

virtio\_net.c : start\_xmit

```
1  static netdev_tx_t start_xmit(struct sk_buff *skb, struct net_device *dev)
2  {
3      struct virtnet_info *vi = netdev_priv(dev);
4      int qnum = skb_get_queue_mapping(skb);
5      struct send_queue *sq = &vi->sq[qnum]; //对应的send_queue
6      int err;
7      struct netdev_queue *txq = netdev_get_tx_queue(dev, qnum);
8      bool kick = !skb->xmit_more;
9
10     /* Free up any pending old buffers before queueing new ones. */
11     /*释放backend处理过的desc*/
12     free_old_xmit_skbs(sq);
13
14     /* Try to transmit */
15     /*发包*/
16     err = xmit_skb(sq, skb);
17
18     /* This should not happen! */
19     if (unlikely(err)) {
20         dev->stats.tx_fifo_errors++;
21         if (net_ratelimit())
22             dev_warn(&dev->dev,
23                     "Unexpected TXQ (%d) queue failure: %d\n", qnum, err);
24         dev->stats.tx_dropped++;
25         kfree_skb(skb);
26         return NETDEV_TX_OK;
27     }
28
29     /* Don't wait up for transmitted skbs to be freed. */
30     /*孤立skb,因为此时对guest来说此skb的内容已发送*/
31     skb_orphan(skb);
32     nf_reset(skb);
33
34     /* Apparently nice girls don't return TX_BUSY; stop the queue
35      * before it gets out of hand. Naturally, this wastes entries. */
36     if (sq->vq->num_free < 2+MAX_SKB_FRAGS) {
37         netif_stop_subqueue(dev, qnum);
38         if (unlikely(!virtqueue_enable_cb_delayed(sq->vq))) {
39             /* More just got used, free them then recheck. */
```

```

40         free_old_xmit_skbs(sq);
41         if (sq->vq->num_free >= 2+MAX_SKB_FRAGS) {
42             netif_start_subqueue(dev, qnum);
43             virtqueue_disable_cb(sq->vq);
44         }
45     }
46 }
47
48     if (kick || netif_xmit_stopped(txq))
49 /*通知host端*/
50         virtqueue_kick(sq->vq);
51
52     return NETDEV_TX_OK;
53 }

```

下面我们来分析**xmit\_skb**函数是如何发包的，主要就是将数据包头部和数据包填入scatterlist，然后调用**virtqueue\_add\_outbuf**将sg table 写入desc描述符表，并将head desc信息写入vring.avail。

virtio\_net.c : xmit\_skb

```

1  static int xmit_skb(struct send_queue *sq, struct sk_buff *skb)
2  {
3      struct virtio_net_hdr_mrg_rxbuf *hdr; //数据包头部
4      const unsigned char *dest = ((struct ethhdr *)skb->data)->h_dest;
5      struct virtnet_info *vi = sq->vq->vdev->priv;
6      unsigned num_sg;
7      unsigned hdr_len = vi->hdr_len;    //Packet virtio header size
8      bool can_push;
9
10     pr_debug("%s: xmit %p %pM\n", vi->dev->name, skb, dest);
11
12     can_push = vi->any_header_sg &&
13         !((unsigned long)skb->data & (__alignof__(*hdr) - 1)) &&
14         !skb_header_cloned(skb) && skb_headroom(skb) >= hdr_len;
15     /* Even if we can, don't push here yet as this would skew
16      * csum_start offset below. */
17     if (can_push)
18         hdr = (struct virtio_net_hdr_mrg_rxbuf *) (skb->data - hdr_len);
19     else
20         hdr = skb_vnet_hdr(skb);
21     /*校验和相关设置*/
22     if (skb->ip_summed == CHECKSUM_PARTIAL) {
23         hdr->hdr.flags = VIRTIO_NET_HDR_F_NEEDS_CSUM;
24         ↑ 90%   hdr->hdr.csum_start = cpu_to_virtio16(vi->vdev,

```

```

25         skb_checksum_start_offset(skb));
26     hdr->hdr.csum_offset = cpu_to_virtio16(vi->vdev,
27         skb->csum_offset);
28 } else {
29     hdr->hdr.flags = 0;
30     hdr->hdr.csum_offset = hdr->hdr.csum_start = 0;
31 }
32 /*如果是GSO数据包需要重新设置包头*/
33     if (skb_is_gso(skb)) {
34         hdr->hdr.hdr_len = cpu_to_virtio16(vi->vdev, skb_headlen(skb));
35         hdr->hdr.gso_size = cpu_to_virtio16(vi->vdev,
36             skb_shinfo(skb)->gso_size);
37         if (skb_shinfo(skb)->gso_type & SKB_GSO_TCPV4)
38             hdr->hdr.gso_type = VIRTIO_NET_HDR_GSO_TCPV4;
39         else if (skb_shinfo(skb)->gso_type & SKB_GSO_TCPV6)
40             hdr->hdr.gso_type = VIRTIO_NET_HDR_GSO_TCPV6;
41         else if (skb_shinfo(skb)->gso_type & SKB_GSO_UDP)
42             hdr->hdr.gso_type = VIRTIO_NET_HDR_GSO_UDP;
43         else
44             BUG();
45         if (skb_shinfo(skb)->gso_type & SKB_GSO_TCP_ECN)
46             hdr->hdr.gso_type |= VIRTIO_NET_HDR_GSO_ECN;
47     } else {
48         hdr->hdr.gso_type = VIRTIO_NET_HDR_GSO_NONE;
49         hdr->hdr.gso_size = hdr->hdr.hdr_len = 0;
50     }
51 /*判断设备是否支持合并buffer*/
52     if (vi->mergeable_rx_bufs)
53         hdr->num_buffers = 0;
54     sg_init_table(sq->sg, MAX_SKB_FRAGS + 2);
55     if (can_push) {
56         __skb_push(skb, hdr_len);
57         num_sg = skb_to_sgvec(skb, sq->sg, 0, skb->len);
58         /* Pull header back to avoid skew in tx bytes calculations. */
59         __skb_pull(skb, hdr_len);
60     } else {
61 /*数据包头部填入scatterlist*/
62         sg_set_buf(sq->sg, hdr, hdr_len);
63 /*数据包填入scatterlist*/
64         num_sg = skb_to_sgvec(skb, sq->sg + 1, 0, skb->len) + 1;
65     }
66 /*sg table 写入desc描述符表, head desc信息写vring.avail*/
67     return virtqueue_add_outbuf(sq->vq, sq->sg, num_sg, skb, GFP_ATOMIC);
68 }

```

可以看到 `sg_set_buf` 和 `skb_to_sgvec` 最后都调用了 `sg_set_page`，`sg_set_page` 中把 `sk_buffer`(逻辑buffer)中的物理块的page信息、offset、len信息放入sg中。

scatterlist.h : `sg_set_buf`

```
1 static inline void sg_set_buf(struct scatterlist *sg, const void *buf,
2                               unsigned int buflen)
3 {
4     #ifdef CONFIG_DEBUG_SG
5         BUG_ON(!virt_addr_valid(buf));
6     #endif
7     sg_set_page(sg, virt_to_page(buf), buflen, offset_in_page(buf));
8 }
```

### `skb_to_sgvec`

skbuff.c : `skb_to_sgvec`

```
1 int skb_to_sgvec(struct sk_buff *skb, struct scatterlist *sg, int offset, int len)
2 {
3     int nsg = __skb_to_sgvec(skb, sg, offset, len);
4
5     sg_mark_end(&sg[nsg - 1]);
6
7     return nsg;
8 }
```

skbuff.c : `__skb_to_sgvec`

```
1 static int
2 __skb_to_sgvec(struct sk_buff *skb, struct scatterlist *sg, int offset, int len)
3 {
4     int start = skb_headlen(skb);
5     int i, copy = start - offset;
6     struct sk_buff *frag_iter;
7     int elt = 0;
8
9     if (copy > 0) {
10         if (copy > len)
11             copy = len;
12         sg_set_buf(sg, skb->data + offset, copy);
13         elt++;
14         if ((len -= copy) == 0)
15             return elt;
```

```

16         offset += copy;
17     }
18
19     for (i = 0; i < skb_shinfo(skb)->nr_frags; i++) {
20         int end;
21
22         WARN_ON(start > offset + len);
23
24         end = start + skb_frag_size(&skb_shinfo(skb)->frags[i]);
25         if ((copy = end - offset) > 0) {
26             skb_frag_t *frag = &skb_shinfo(skb)->frags[i];
27
28             if (copy > len)
29                 copy = len;
30             sg_set_page(&sg[elt], skb_frag_page(frag), copy,
31                 frag->page_offset+offset-start);
32             elt++;
33             if (!(len -= copy))
34                 return elt;
35             offset += copy;
36         }
37         start = end;
38     }
39
40     skb_walk_frags(skb, frag_iter) {
41         int end;
42
43         WARN_ON(start > offset + len);
44
45         end = start + frag_iter->len;
46         if ((copy = end - offset) > 0) {
47             if (copy > len)
48                 copy = len;
49             elt += __skb_to_sgvec(frag_iter, sg+elt, offset - start,
50                 copy);
51             if ((len -= copy) == 0)
52                 return elt;
53             offset += copy;
54         }
55         start = end;
56     }
57     BUG_ON(len);
58     return elt;
59 }

```

```

1 static inline void sg_set_page(struct scatterlist *sg, struct page *page,
2                               unsigned int len, unsigned int offset)
3 {
4     sg_assign_page(sg, page);
5     sg->offset = offset;
6     sg->length = len;
7 }

```

下面看**virtqueue\_add\_outbuf**如何将sg table 写入desc描述符表。

```

1 int virtqueue_add_outbuf(struct virtqueue *vq,
2                          struct scatterlist *sg, unsigned int num,
3                          void *data,
4                          gfp_t gfp)
5 {
6     return virtqueue_add(vq, &sg, num, 1, 0, data, gfp);
7 }

```

**virtqueue\_add**首先判断是否支持间接描述符表并且 $total\_sg > 1$ (否则也不需要使用间接描述符表了), 如果支持indirect把descs\_used设置为1, 否则设为实际的sg entry总数total\_sg。如果主描述符表的空闲表项数小于total\_sg就错误返回。然后在for循环中先调用**vring\_map\_one\_sg**函数得到scatterlist中数据的GPA,然后将sg table中entry信息分别对应到描述符表中。标记last desc信息, 更新空闲desc数量以及free\_head, 最后将head desc信息写入vring.avail->ring[]并更新idx。

```

1 static inline int virtqueue_add(struct virtqueue *_vq,
2                                struct scatterlist *sgs[],
3                                unsigned int total_sg,
4                                unsigned int out_sgs,
5                                unsigned int in_sgs,
6                                void *data,
7                                gfp_t gfp)
8 {
9     struct vring_virtqueue *vq = to_vvq(_vq);
10    struct scatterlist *sg;
11    struct vring_desc *desc;
12    unsigned int i, n, avail, descs_used, uninitialized_var(prev), err_idx;
13    int head;
14    bool indirect;
15
16    START_USE(vq);
17    .....
18    head = vq->free_head;

```

```

19 //判断是否支持间接描述符并且total_sg > 1
20 /* If the host supports indirect descriptor tables, and we have multiple
21  * buffers, then go indirect. FIXME: tune this threshold */
22 if (vq->indirect && total_sg > 1 && vq->vq.num_free)
23     desc = alloc_indirect(_vq, total_sg, gfp);
24 else
25     desc = NULL;
26
27 if (desc) {
28     /* Use a single buffer which doesn't continue */
29     indirect = true;
30     /* Set up rest to use this indirect table. */
31     i = 0;
32     desc_used = 1;
33 } else {
34     indirect = false;
35     desc = vq->vring.desc;
36     i = head;
37     desc_used = total_sg;
38 }
39 /*主描述符表的空闲表项数小于total_sg，错误*/
40 if (vq->vq.num_free < desc_used) {
41     .....
42     return -ENOSPC;
43 }
44 /*sg table中entry信息记录到对应的desc表中*/
45 for (n = 0; n < out_sgs; n++) {
46     for (sg = sgs[n]; sg; sg = sg_next(sg)) {
47         /*得到总线地址addr，也就是GPA*/
48         dma_addr_t addr = vring_map_one_sg(vq, sg, DMA_TO_DEVICE);
49         if (vring_mapping_error(vq, addr))
50             goto unmap_release;
51         /*GPA等信息填入desc[]中*/
52         desc[i].flags = cpu_to_virtio16(_vq->vdev, VRING_DESC_F_NEXT);
53         desc[i].addr = cpu_to_virtio64(_vq->vdev, addr);
54         desc[i].len = cpu_to_virtio32(_vq->vdev, sg->length);
55         prev = i;
56         i = virtio16_to_cpu(_vq->vdev, desc[i].next);
57     }
58 }
59 for (; n < (out_sgs + in_sgs); n++) {
60     for (sg = sgs[n]; sg; sg = sg_next(sg)) {
61         dma_addr_t addr = vring_map_one_sg(vq, sg, DMA_FROM_DEVICE);
62         if (vring_mapping_error(vq, addr))
63             goto unmap_release;
64
65         desc[i].flags = cpu_to_virtio16(_vq->vdev, VRING_DESC_F_NEXT | VRING_DESC_F_WRITE);

```



```

66         desc[i].addr = cpu_to_virtio64(_vq->vdev, addr);
67         desc[i].len = cpu_to_virtio32(_vq->vdev, sg->length);
68         prev = i;
69         i = virtio16_to_cpu(_vq->vdev, desc[i].next);
70     }
71 }
72 /*标记last desc信息*/
73 /* Last one doesn't continue. */
74 desc[prev].flags &= cpu_to_virtio16(_vq->vdev, ~VRING_DESC_F_NEXT);
75
76 if (indirect) {
77     /* Now that the indirect table is filled in, map it. */
78     dma_addr_t addr = vring_map_single(
79         vq, desc, total_sg * sizeof(struct vring_desc),
80         DMA_TO_DEVICE);
81     if (vring_mapping_error(vq, addr))
82         goto unmap_release;
83     vq->vring.desc[head].flags = cpu_to_virtio16(_vq->vdev, VRING_DESC_F_INDIRECT);
84     vq->vring.desc[head].addr = cpu_to_virtio64(_vq->vdev, addr);
85     vq->vring.desc[head].len = cpu_to_virtio32(_vq->vdev, total_sg * sizeof(struct vring_desc));
86 }
87 /*更新空闲desc数量*/
88 /* We're using some buffers from the free list. */
89 vq->vq.num_free -= descs_used;
90 /*更新free_head指针*/
91 /* Update free pointer */
92 if (indirect)
93     vq->free_head = virtio16_to_cpu(_vq->vdev, vq->vring.desc[head].next);
94 else
95     vq->free_head = i;
96
97 /* Store token and indirect buffer state. */
98 vq->desc_state[head].data = data;
99 if (indirect)
100     vq->desc_state[head].indir_desc = desc;
101 vq->vq.num_free -= descs_used;
102 /*desc链中第一个desc放入vring_avail中的ring[]*/
103 /* Put entry in available array (but don't update avail->idx until they
104  * do sync). */
105 avail = virtio16_to_cpu(_vq->vdev, vq->vring.avail->idx) & (vq->vring.num - 1);
106 vq->vring.avail->ring[avail] = cpu_to_virtio16(_vq->vdev, head);
107
108 /* Descriptors and available array need to be set before we expose the
109  * new available array entries. */
110 virtio_wmb(vq->weak_barriers);
111 /*更新vring_avail中idx*/
112 vq->vring.avail->idx = cpu_to_virtio16(_vq->vdev, virtio16_to_cpu(_vq->vdev, vq->vring.avail->i

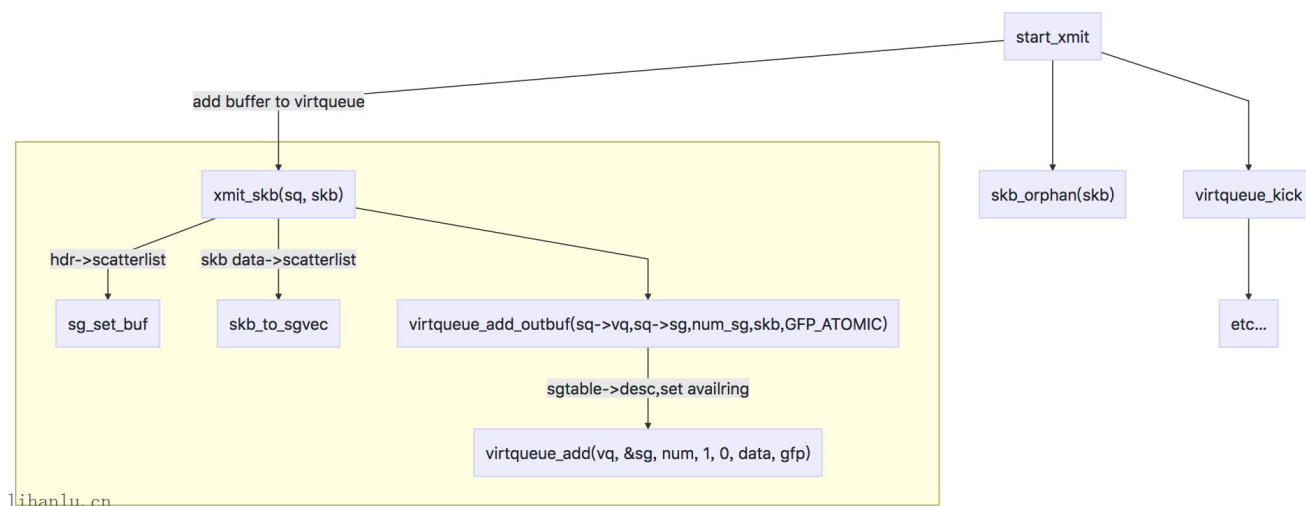
```

```

113     vq->num_added++;
114     .....
115 }

```

回到xmit\_skb之后调用virtqueue\_kick通知Host(具体的通知过程会另写一篇文章分析),上述过程总结如下图。



## Backend

Backend收到通知后会执行virtio\_queue\_host\_notifier\_read函数，我们从这里开始分析

```

1 void virtio_queue_host_notifier_read(EventNotifier *n)
2 {
3     VirtQueue *vq = container_of(n, VirtQueue, host_notifier);
4     if (event_notifier_test_and_clear(n)) {
5         virtio_queue_notify_vq(vq);
6     }
7 }

```

virtio\_queue\_notify\_vq中调用了virtio\_net\_device\_realize中virtio\_net\_add\_queue时绑定的处理函数handle\_output，该函数根据不同的设备有不同的实现，以网卡为例看看创建VirtQueue的时候给绑定的是哪个函数。

```

1 static void virtio_queue_notify_vq(VirtQueue *vq)
2 {
3     if (vq->vring.desc && vq->handle_output) {
4         VirtIODevice *vdev = vq->vdev;
5
6         if (unlikely(vdev->broken)) {
7             return;
8         }
9     }

```

```

10     trace_virtio_queue_notify(vdev, vq - vdev->vq, vq);
11     vq->handle_output(vdev, vq);
12 }
13 }

```

可以看到这里给 rx\_vq 绑定的是 `virtio_net_handle_rx`，tx\_vq 绑定的 `virtio_net_handle_tx_timer`或`virtio_net_handle_tx_bh`，下面以`virtio_net_handle_tx_bh`为例看一下。

```

1  static void virtio_net_add_queue(VirtIONet *n, int index)
2  {
3      VirtIODevice *vdev = VIRTIO_DEVICE(n);
4
5      n->vqs[index].rx_vq = virtio_add_queue(vdev, n->net_conf.rx_queue_size,
6                                              virtio_net_handle_rx);
7
8      if (n->net_conf.tx && !strcmp(n->net_conf.tx, "timer")) {
9          n->vqs[index].tx_vq =
10             virtio_add_queue(vdev, n->net_conf.tx_queue_size,
11                             virtio_net_handle_tx_timer);
12          n->vqs[index].tx_timer = timer_new_ns(QEMU_CLOCK_VIRTUAL,
13                                                virtio_net_tx_timer,
14                                                &n->vqs[index]);
15      } else {
16          n->vqs[index].tx_vq =
17             virtio_add_queue(vdev, n->net_conf.tx_queue_size,
18                             virtio_net_handle_tx_bh);
19          n->vqs[index].tx_bh = qemu_bh_new(virtio_net_tx_bh, &n->vqs[index]);
20      }
21
22      n->vqs[index].tx_waiting = 0;
23      n->vqs[index].n = n;
24 }

```

`virtio_net_handle_tx_bh`中调用了`qemu_bh_schedule(q->tx_bh)`运行参数指定的函数，而参数tx\_bh在`virtio_net_add_queue`中调用`qemu_bh_new(virtio_net_tx_bh, &n->vqs[index])`绑定为`virtio_net_tx_bh`

```

1  static void virtio_net_handle_tx_bh(VirtIODevice *vdev, VirtQueue *vq)
2  {
3      VirtIONet *n = VIRTIO_NET(vdev);
4      VirtIONetQueue *q = &n->vqs[vq2q(virtio_get_queue_index(vq))];
5       if (unlikely((n->status & VIRTIO_NET_S_LINK_UP) == 0)) {
6

```

```

7     virtio_net_drop_tx_queue_data(vdev, vq);
8     return;
9 }
10
11 if (unlikely(q->tx_waiting)) {
12     return;
13 }
14 q->tx_waiting = 1;
15 /* This happens when device was stopped but VCPU wasn't. */
16 if (!vdev->vm_running) {
17     return;
18 }
19 virtio_queue_set_notification(vq, 0);
20 qemu_bh_schedule(q->tx_bh);
21 }

```

### virtio\_net\_tx\_bh函数调用发送函数virtio\_net\_flush\_tx

```

1 static void virtio_net_tx_bh(void *opaque)
2 {
3     .....
4     ret = virtio_net_flush_tx(q);
5     .....
6 }

```

**virtio\_net\_flush\_tx**函数调用**virtqueue\_pop**从avail ring中取数据buffer的head desc，然后调用**qemu\_sendv\_packet\_async**发送packet，**qemu\_sendv\_packet\_async**函数本文不做详细的分析。

```

1 static int32_t virtio_net_flush_tx(VirtIONetQueue *q)
2 {
3     VirtIONet *n = q->n;
4     VirtIODevice *vdev = VIRTIO_DEVICE(n);
5     VirtQueueElement *elem;
6     int32_t num_packets = 0;
7     int queue_index = vq2q(virtio_get_queue_index(q->tx_vq));
8     if (!(vdev->status & VIRTIO_CONFIG_S_DRIVER_OK)) {
9         return num_packets;
10    }
11
12    if (q->async_tx.elem) {
13        virtio_queue_set_notification(q->tx_vq, 0);
14        return num_packets;
15    }
16

```

```

17     for (;;) {
18         ssize_t ret;
19         unsigned int out_num;
20         struct iovec sg[VIRTQUEUE_MAX_SIZE], sg2[VIRTQUEUE_MAX_SIZE + 1], *out_sg;
21         struct virtio_net_hdr_mrg_rxbuf mhdr;
22         //从avail ring中取数据buffer的head desc
23         elem = virtqueue_pop(q->tx_vq, sizeof(VirtQueueElement));
24         if (!elem) {
25             break;
26         }
27
28         out_num = elem->out_num;
29         out_sg = elem->out_sg;
30         if (out_num < 1) {
31             virtio_error(vdev, "virtio-net header not in first element");
32             virtqueue_detach_element(q->tx_vq, elem, 0);
33             g_free(elem);
34             return -EINVAL;
35         }
36
37         if (n->has_vnet_hdr) {
38             if (iov_to_buf(out_sg, out_num, 0, &mhdr, n->guest_hdr_len) <
39                 n->guest_hdr_len) {
40                 virtio_error(vdev, "virtio-net header incorrect");
41                 virtqueue_detach_element(q->tx_vq, elem, 0);
42                 g_free(elem);
43                 return -EINVAL;
44             }
45             if (n->needs_vnet_hdr_swap) {
46                 virtio_net_hdr_swap(vdev, (void *) &mhdr);
47                 sg2[0].iov_base = &mhdr;
48                 sg2[0].iov_len = n->guest_hdr_len;
49                 out_num = iov_copy(&sg2[1], ARRAY_SIZE(sg2) - 1,
50                     out_sg, out_num,
51                     n->guest_hdr_len, -1);
52                 if (out_num == VIRTQUEUE_MAX_SIZE) {
53                     goto drop;
54                 }
55                 out_num += 1;
56                 out_sg = sg2;
57             }
58         }
59         /*
60          * If host wants to see the guest header as is, we can
61          * pass it on unchanged. Otherwise, copy just the parts
62          * that host is interested in.
63         */

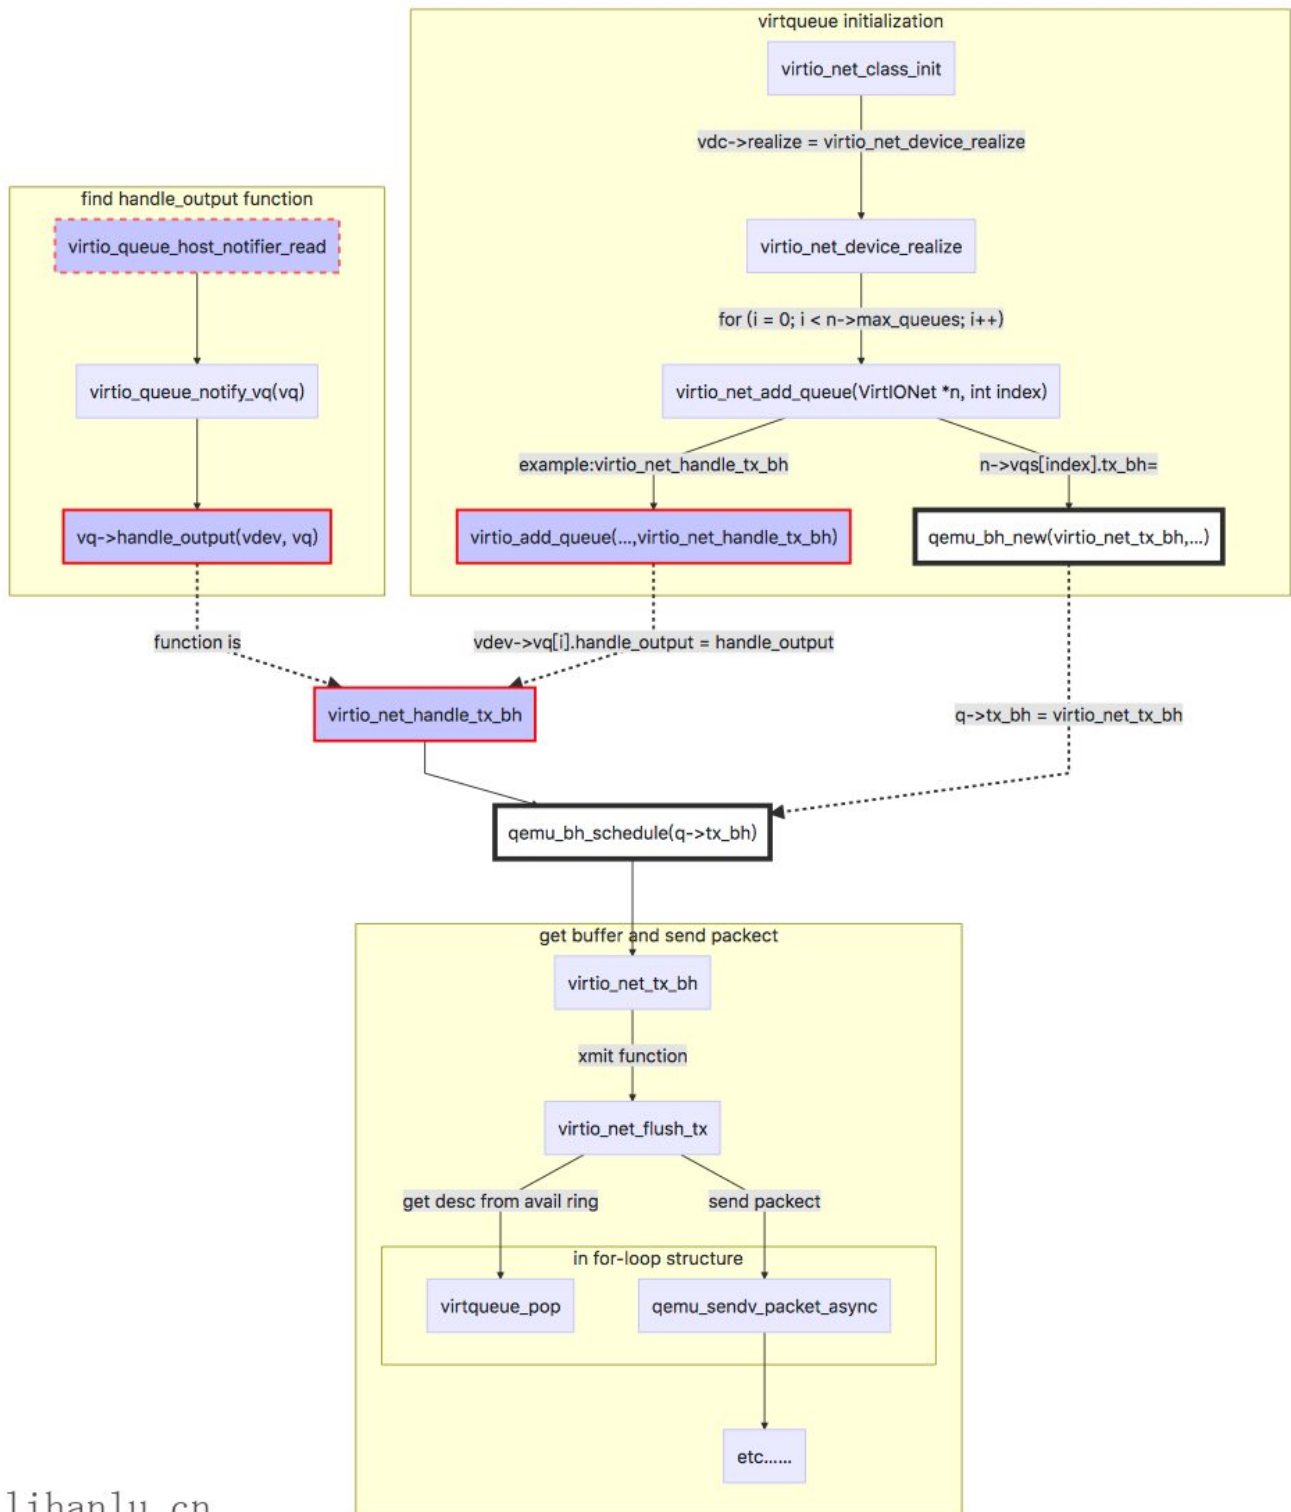
```

```

64     assert(n->host_hdr_len <= n->guest_hdr_len);
65     if (n->host_hdr_len != n->guest_hdr_len) {
66         unsigned sg_num = iov_copy(sg, ARRAY_SIZE(sg),
67                                     out_sg, out_num,
68                                     0, n->host_hdr_len);
69         sg_num += iov_copy(sg + sg_num, ARRAY_SIZE(sg) - sg_num,
70                             out_sg, out_num,
71                             n->guest_hdr_len, -1);
72         out_num = sg_num;
73         out_sg = sg;
74     }
75     /*调用qemu发包函数，通过qemu nic发送*/
76     ret = qemu_sendv_packet_async(qemu_get_subqueue(n->nic, queue_index),
77                                   out_sg, out_num, virtio_net_tx_complete);
78     if (ret == 0) {
79         virtio_queue_set_notification(q->tx_vq, 0);
80         q->async_tx.elem = elem;
81         return -EBUSY;
82     }
83
84 drop:
85     /*取消内存映射，更新used vring信息*/
86     virtqueue_push(q->tx_vq, elem, 0);
87     virtio_notify(vdev, q->tx_vq);
88     g_free(elem);
89
90     if (++num_packets >= n->tx_burst) {
91         break;
92     }
93 }
94 return num_packets;
95 }

```

Host收到通知后的发包流程总结如下图。



lihanlu.cn

## 联系我

你可以直接在下方留言，也可以✉E-Mail联系我。

打赏

本文链接: <http://lihanlu.cn/virtio-net-xmit/>

版权声明: 本博客所有文章除特别声明外, 均采用 [CC BY-NC-SA](#) 许可协议。转载请注明出处!

virtio

< Virtio原理简介

virtio前端通知机制分析 >

昵称	邮箱	网址(http://)
Just go go		
<div></div>		
		<div></div> <div>提交</div>

## 1 评论



Anonymous

Chrome 80.0.3987.132 Windows 10.0

2020-03-19

回复

先读一下,有问题再请教

Powered By Valine  
v1.4.4

京ICP备19012335号-1

© 2017 - 2020  Lauren |  92k |  1:23

↑ 90%

由 [Hexo](#) & [NexT.Mist](#) 强力驱动



